

# The Sleeping Teaching Assistant

A Synchronization Problem using POSIX Threads, Mutexes & Semaphores

---

## Team Members

Mohammad Ikhlas – AP23110010087

Praveen Vemulapati – AP23110011455


Venkatesh Manipati – AP23110011626

# Problem Statement

A university TA helps students in an office with one desk and three waiting chairs in the hallway.

- If no students are present, the TA naps.
- If a student arrives and the TA is sleeping, the student must awaken the TA.
- If a student arrives and the TA is busy, the student waits in a hallway chair.
- If all hallway chairs are full, the student leaves and will come back later.

**Physical Analogy:** This is a classic "Producer-Consumer" problem. Students "produce" requests for help, and the TA "consumes" them.



# Goals & Constraints

⌚ **Goal:** Model this scenario concurrently, ensuring the TA and Students operate correctly without errors (race conditions or deadlocks).

👤 **Constraint: 1 TA (Single Consumer)**

Only one "help session" can be active at a time. The TA must process requests one by one.

🪑 **Constraint: 3 Chairs (Bounded Buffer)**

The waiting area has a fixed size. We can't have more than 3 students waiting.


🏃 **Constraint: Leaving Policy ("Balking")**

If the buffer (hallway) is full, arriving students (producers) do not wait; they leave immediately.

# Design & Flow

## Logical Flow (Student)

1. Student thread is created (arrives).
2. Attempt to acquire lock on chairs.
3. Check `waiting\_count`:
  - If `count < 3` (Space): Increment count, signal TA (wake up), wait for TA to be free, then get help.
  - If `count == 3` (Full): Release lock and leave.
4. Thread exits.



# Synchronization Primitives



## **pthread\_t (Threads)**

We use POSIX threads to model concurrency. The TA is one long-running thread, and each student is a new thread that is created and joins.



## **Mutex (The "Talking Stick")**

A `pthread\_mutex\_t` is used for \*\*Mutual Exclusion\*\*. It protects the shared `waiting\_count` variable. Only the thread holding the lock can read or write this count, preventing race conditions.



## **Semaphores (The "Signals")**

`sem\_t` is used for \*\*Signaling\*\*.  
We use two:  
**sem\_students (Doorbell):**  
Counts waiting students. TA sleeps on this.  
**sem\_ta (Office Door):** Signals the TA is ready. Students wait on this.

# Algorithm Pseudocode

## Student (Loop)

```
Student_Thread() { lock(mutex); if (waiting_count < 3) { waiting_count++; unlock(mutex); // Ring doorbell to wake TA post(sem_student); sleep(); } }
```

# Implementation: The "Handshake"

## Key Code Snippets

```
// --- STUDENT --- // Wake up the TA sem_post(&sem_students); // Wait for TA to be ready sem_wait(&sem_ta); // --- TA --- // Wait for
```

# Correctness Reasoning

- 🛡 **Race Conditions Prevented:** All access to the shared `waiting\_count` variable is "mutually exclusive," protected by the `mutex`. Only one thread can read or write it at a time.
- 🔔 **"Lost Wakeups" Avoided:** Semaphores are counters, not flags. If a student signals (`sem\_post`) \*before\* the TA sleeps (`sem\_wait`), the semaphore count becomes 1. When the TA later calls `sem\_wait`, it sees the count is 1, decrements it, and \*\*does not sleep\*\*.
- ⌚ **Deadlock Avoided:** Deadlock ("Circular Wait") is impossible. A student \*never\* holds the mutex while waiting for a semaphore. They `unlock(mutex)` \*before\* they `wait(sem\_ta)`, breaking the chain.

Students arrive one at a time. Tests the TA's sleep/wake cycle.

- **Test 2:**  
**Burst Arrival**  
5 students arrive at once. Tests the "balking" logic. The first 3 should wait, the next 2 should leave.

- **Edge Case:**  
Student arrives

# Extensions & Improvements

## Multiple TAs

The project could be extended to support multiple TA threads. This would change the logic significantly:

- We would need a semaphore to count \*available TAs\*.
- Students would `wait()` on this "available TA" semaphore.
- This turns the problem into a more complex "Multi-Producer, Multi-Consumer" model.

## Student Priorities

We could implement a system where some students get priority (e.g., final-year vs. first-year).

- This would require replacing the simple `waiting\_count` with a \*\*Priority Queue\*\* data structure.
- The \*same mutex\* would be used to protect the queue, but the TA would "dequeue" the highest-priority student.

# Conclusion & Lessons Learned



## The Right Tool for the Job

This project highlights the difference between synchronization tools.

**Mutex:** Use for **Protection** (protecting shared data like a variable).

**Semaphore:** Use for **Coordination** (signaling, sleeping, and waking threads).



## Concurrency is Manageable

Handling concurrency pitfalls like race conditions and deadlocks is complex, but possible with a careful design. By strictly enforcing lock order and using semaphores correctly, we can build a robust, efficient, and correct system.

# References

-  Silberschatz, Galvin, & Gagne. (2018). *Operating System Concepts (10th Ed.)*.
-  Downey, Allen B. (2016). *The Little Book of Semaphores*.
-  POSIX Threads Programming Tutorial (GeeksforGeeks, IBM Developer).

# Questions?


Thank you.

# Image Sources




<https://media.istockphoto.com/id/1158203930/vector/burnout-concept-illustration-with-exhausted-man-office-worker-sitting-at-the-table.jpg?r=612x612&w=0&k=20&c=jleYMmABeH5fjh7wc--ySzdPmyHml2Lyuy3GmcNdzhE=>

Source: [www.istockphoto.com](http://www.istockphoto.com)



[https://miro.medium.com/v2/resize:fit:846/1\\*JD93hAvRCrR7uxyDVutRKw.png](https://miro.medium.com/v2/resize:fit:846/1*JD93hAvRCrR7uxyDVutRKw.png)

Source: [levelup.gitconnected.com](http://levelup.gitconnected.com)



<https://i.stack.imgur.com/e92Ju.png>

Source: [stackoverflow.com](http://stackoverflow.com)